

Multiple MLP Architectures on MNIST

December 8, 2018

1 Multiple MLP Architectures using Keras on MNIST

1.1 Purpose

The purpose of the study is to try out **3 different MLP architectures on MNIST dataset to compare the performance**. The implementation is done in Keras.

1.2 Steps at a Glance:

1. Take the famous MNIST dataset as input. <http://yann.lecun.com/exdb/mnist/>
2. Feed it into **2-layered MLP Architecture: Input(784)-ReLu(512)-ReLu(128)-Sigmoid(output)**
3. Find the accuracy and draw the **Loss vs Epoch Plot**
4. Introduce **Batch Normalization and Dropouts**.
5. Evaluate the model again by estimating accuracy and drawing loss diagram.
6. Feed same input to **3 layered MLP Architecture: Input(784)-ReLu(512)-ReLu(256)-ReLu(64)-Sigmoid(output)**
7. Introduce Batch Normalization and Dropouts & evaluate the model again.
8. Feed same input to **5 layered MLP Architecture: Input(784)-ReLu(512)-ReLu(256)-ReLu(144)-ReLu(96)-ReLu(36)-Sigmoid(output)**
9. Introduce Batch Normalization and Dropouts & evaluate the model again.
10. Analyze the output from the above 3 architectures and draw conclusions .

```
In [53]: # if you keras is not using tensorflow as backend
        # set "KERAS_BACKEND=tensorflow" use this command
        from keras.utils import np_utils
        from keras.datasets import mnist
        import seaborn as sns
        from keras.initializers import RandomNormal
```

1.3 Loading and Pre-processing Data

```
In [54]: # the data, shuffled and split between train and test sets
        (X_train, y_train), (X_test, y_test) = mnist.load_data()

In [55]: print("Number of training examples :", X_train.shape[0],
              "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
        print("Number of testing examples :", X_test.shape[0],
              "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))
```

Number of training examples : 60000 and each image is of shape (28, 28)
 Number of testing examples : 10000 and each image is of shape (28, 28)

```
In [56]: # if you observe the input shape its 3 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784
```

```
X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])
```

```
In [57]: # after converting the input images from 3d to 2d vectors
```

```
print("Number of training examples :", X_train.shape[0],
      "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of testing examples :", X_test.shape[0],
      "and each image is of shape (%d)"%(X_test.shape[1]))
```

Number of training examples : 60000 and each image is of shape (784)
 Number of testing examples : 10000 and each image is of shape (784)

```
In [58]: # An example data point
print(X_train[0])
```

```
[ 0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175  26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30  36  94 154
170 253 253 253 253 253 225 172 253 242 195  64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251  93  82
 82  56  39  0  0  0  0  0  0  0  0  0  0  0  0  18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  80 156 107 253 253 205  11  0  43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253  90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  139 253 190  2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253  70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  81 240 253 253 119  25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0]
```

```

0 0 45 186 253 253 150 27 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 16 93 252 253 187
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 249 253 249 64 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 46 130 183 253
253 207 2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 39 148 229 253 253 253 250 182 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 24 114 221 253 253 253
253 201 78 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 23 66 213 253 253 253 253 198 81 2 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 18 171 219 253 253 253 253 195
80 9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
55 172 226 253 253 253 253 244 133 11 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 136 253 253 253 212 135 132 16
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
0 0 0 0 0 0 0 0 0 0 0]

```

```

In [59]: # if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
#  $X \Rightarrow (X - X_{min}) / (X_{max} - X_{min}) = X / 255$ 

```

```

X_train = X_train/255
X_test = X_test/255

```

```

In [60]: # example data point after normlizing
print(X_train[0])

```

```

[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.]

```

| | | | | | |
|------------|------------|------------|------------|------------|------------|
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.01176471 | 0.07058824 | 0.07058824 | 0.07058824 |
| 0.49411765 | 0.53333333 | 0.68627451 | 0.10196078 | 0.65098039 | 1. |
| 0.96862745 | 0.49803922 | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.11764706 | 0.14117647 | 0.36862745 | 0.60392157 |
| 0.66666667 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.88235294 | 0.6745098 | 0.99215686 | 0.94901961 | 0.76470588 | 0.25098039 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0.19215686 |
| 0.93333333 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.98431373 | 0.36470588 | 0.32156863 |
| 0.32156863 | 0.21960784 | 0.15294118 | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0.07058824 | 0.85882353 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.77647059 | 0.71372549 |
| 0.96862745 | 0.94509804 | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.31372549 | 0.61176471 | 0.41960784 | 0.99215686 |
| 0.99215686 | 0.80392157 | 0.04313725 | 0. | 0.16862745 | 0.60392157 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0.05490196 | 0.00392157 | 0.60392157 | 0.99215686 | 0.35294118 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0.54509804 | 0.99215686 | 0.74509804 | 0.00784314 | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0.04313725 |
| 0.74509804 | 0.99215686 | 0.2745098 | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0.1372549 | 0.94509804 |
| 0.88235294 | 0.62745098 | 0.42352941 | 0.00392157 | 0. | 0. |

| | | | | | |
|------------|------------|------------|------------|------------|------------|
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0.31764706 | 0.94117647 | 0.99215686 |
| 0.99215686 | 0.46666667 | 0.09803922 | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.17647059 | 0.72941176 | 0.99215686 | 0.99215686 |
| 0.58823529 | 0.10588235 | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0.0627451 | 0.36470588 | 0.98823529 | 0.99215686 | 0.73333333 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0.97647059 | 0.99215686 | 0.97647059 | 0.25098039 | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.18039216 | 0.50980392 | 0.71764706 | 0.99215686 |
| 0.99215686 | 0.81176471 | 0.00784314 | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0.15294118 | 0.58039216 |
| 0.89803922 | 0.99215686 | 0.99215686 | 0.99215686 | 0.98039216 | 0.71372549 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0.09411765 | 0.44705882 | 0.86666667 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.78823529 | 0.30588235 | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0.09019608 | 0.25882353 | 0.83529412 | 0.99215686 |
| 0.99215686 | 0.99215686 | 0.99215686 | 0.77647059 | 0.31764706 | 0.00784314 |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0.07058824 | 0.67058824 |
| 0.85882353 | 0.99215686 | 0.99215686 | 0.99215686 | 0.99215686 | 0.76470588 |
| 0.31372549 | 0.03529412 | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0.21568627 | 0.6745098 | 0.88627451 | 0.99215686 | 0.99215686 | 0.99215686 |
| 0.99215686 | 0.95686275 | 0.52156863 | 0.04313725 | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |
| 0. | 0. | 0. | 0. | 0. | 0. |

[illegible]

```
In [61]: # here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ", Y_train[0])
```

Class label of first image : 5
After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```
In [62]: # https://keras.io/getting-started/sequential-model-guide/

# The Sequential model is a linear stack of layers.
# you can create a Sequential model by passing
# a list of layer instances to the constructor:

# model = Sequential([
#     Dense(32, input_shape=(784,)),
#     Activation('relu'),
#     Dense(10),
#     Activation('softmax'),
```

```

# ])

# You can also simply add layers via the .add() method:

# model = Sequential()
# model.add(Dense(32, input_dim=784))
# model.add(Activation('relu'))

###

# https://keras.io/layers/core/

# keras.layers.Dense(units, activation=None, use_bias=True,
# kernel_initializer='glorot_uniform', bias_initializer='zeros',
# kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,
# kernel_constraint=None, bias_constraint=None)

# Dense implements the operation: output = activation(dot(input, kernel) + bias)
# where activation is the element-wise activation function passed as the activation
# argument, kernel is a weights matrix created by the layer, and
# bias is a bias vector created by the layer (only applicable if use_bias is True).

# output = activation(dot(input, kernel) + bias) => y = activation(WT. X + b)

####

# https://keras.io/activations/

# Activations can either be used through an Activation layer, or through
# the activation argument supported by all forward layers:

# from keras.layers import Activation, Dense

# model.add(Dense(64))
# model.add(Activation('tanh'))

# This is equivalent to:
# model.add(Dense(64, activation='tanh'))

# there are many activation functions available ex: tanh, relu, softmax

from keras.models import Sequential
from keras.layers import Dense, Activation

```

2 Initializations

In [63]: *# initialization of some model parameters*

```
output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20
```

3 Custom-Defined Functions

In [64]: *%matplotlib inline*

```
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(fig, x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Validation Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()
```

In [65]: *# To train the model using Adam*

```
# This function is common to all models.
def trainModel(model):
    model.compile(optimizer='adam',
                  loss='categorical_crossentropy', metrics=['accuracy'])

    history = model.fit(X_train, Y_train, batch_size=batch_size,
                        epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

    return history
```

In [66]: *# To plot the Train & Test loss graph.*

```
# This function is common to all models.
def plotGraph(model, history):
    score = model.evaluate(X_test, Y_test, verbose=0)
    print('Test score:', score[0])
    print('Test accuracy:', score[1])

    fig, ax = plt.subplots(1, 1)
    ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

    # list of epoch numbers
```



```

x = list(range(1,nb_epoch+1))

# print(history.history.keys())
# dict_keys(['val_loss', 'val_acc', 'loss', 'acc'])
# history = model_drop.fit(X_train, Y_train, batch_size=batch_size,
#                           #epochs=nb_epoch, verbose=1, validation_data=(X_test, Y_test))

# we will get val_loss and val_acc only
# when you pass the paramter validation_data
# val_loss : validation loss
# val_acc : validation accuracy

# loss : training loss
# acc : train accuracy
# for each key in history.history we will have
# a list of length equal to number of epochs

vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(fig, x, vy, ty, ax)

```

3.1 Model 1: (2-layered MLP Architecture)

```

In [67]: def plotWeightM1(model):
    w_after = model.get_weights()

    h1_w = w_after[0].flatten().reshape(-1,1)
    h2_w = w_after[2].flatten().reshape(-1,1)
    out_w = w_after[4].flatten().reshape(-1,1)

    fig = plt.figure()
    plt.title("Weight matrices after model trained")
    plt.subplot(1, 3, 1)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=h1_w,color='b')
    plt.xlabel('Hidden Layer 1')

    plt.subplot(1, 3, 2)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=h2_w, color='r')
    plt.xlabel('Hidden Layer 2 ')

    plt.subplot(1, 3, 3)
    plt.title("Trained Weights")
    ax = sns.violinplot(y=out_w,color='y')
    plt.xlabel('Output Layer ')
    plt.show()

```

3.1.1 Input(784)-ReLu(512)-ReLu(128)-Softmax(output)


```
In [68]: model_relu = Sequential()
        model_relu.add(Dense(512, activation='relu',
                             input_shape=(input_dim,), kernel_initializer='he_normal'))
        model_relu.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(output_dim, activation='softmax'))


        print(model_relu.summary())

        history = trainModel(model=model_relu)
        plotGraph(model=model_relu, history=history)
        plotWeightM1(model=model_relu)
```

```
-----
Layer (type)                 Output Shape              Param #
=====
dense_79 (Dense)             (None, 512)              401920
-----
dense_80 (Dense)             (None, 128)              65664
-----
dense_81 (Dense)             (None, 10)               1290
=====
Total params: 468,874
Trainable params: 468,874
Non-trainable params: 0
-----
None
Train on 60000 samples, validate on 10000 samples
Epoch 1/20
60000/60000 [=====] - 16s 268us/step - loss: 0.2330 - acc: 0.9315 - val
Epoch 2/20
60000/60000 [=====] - 4s 62us/step - loss: 0.0859 - acc: 0.9741 - val
Epoch 3/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0531 - acc: 0.9836 - val
Epoch 4/20
60000/60000 [=====] - 3s 55us/step - loss: 0.0380 - acc: 0.9883 - val
Epoch 5/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0280 - acc: 0.9910 - val
Epoch 6/20
60000/60000 [=====] - 3s 57us/step - loss: 0.0223 - acc: 0.9929 - val
Epoch 7/20
60000/60000 [=====] - 3s 58us/step - loss: 0.0179 - acc: 0.9942 - val
Epoch 8/20
60000/60000 [=====] - 4s 60us/step - loss: 0.0130 - acc: 0.9959 - val
Epoch 9/20
60000/60000 [=====] - 4s 59us/step - loss: 0.0131 - acc: 0.9956 - val
Epoch 10/20
60000/60000 [=====] - 4s 58us/step - loss: 0.0120 - acc: 0.9963 - val
```

Epoch 11/20
 60000/60000 [=====] - 3s 57us/step - loss: 0.0139 - acc: 0.9956 - val.
 Epoch 12/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0096 - acc: 0.9966 - val.
 Epoch 13/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0085 - acc: 0.9972 - val.
 Epoch 14/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0113 - acc: 0.9964 - val.
 Epoch 15/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0079 - acc: 0.9973 - val.
 Epoch 16/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0068 - acc: 0.9979 - val.
 Epoch 17/20
 60000/60000 [=====] - 3s 57us/step - loss: 0.0053 - acc: 0.9986 - val.
 Epoch 18/20
 60000/60000 [=====] - 3s 57us/step - loss: 0.0041 - acc: 0.9988 - val.
 Epoch 19/20
 60000/60000 [=====] - 3s 58us/step - loss: 0.0091 - acc: 0.9971 - val.
 Epoch 20/20
 60000/60000 [=====] - 3s 56us/step - loss: 0.0114 - acc: 0.9964 - val.
 Test score: 0.096843903848933
 Test accuracy: 0.9812





3.1.2 Model 1: M1 + Batch-Normalization on hidden Layers

In [69]: *# Multilayer perceptron*

```
# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution N(0,)
# we satisfy this condition with  $= (2/(n_i+n_{i+1}))$ .
# h1 =>  $= (2/(n_i+n_{i+1})) = 0.039 \Rightarrow N(0, ) = N(0, 0.039)$ 
# h2 =>  $= (2/(n_i+n_{i+1})) = 0.055 \Rightarrow N(0, ) = N(0, 0.055)$ 
# h1 =>  $= (2/(n_i+n_{i+1})) = 0.120 \Rightarrow N(0, ) = N(0, 0.120)$ 

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='relu',
                      input_shape=(input_dim,), kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())

model_batch.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

```
history = trainModel(model=model_batch)
plotGraph(model=model_batch, history=history)
plotWeightM1(model=model_batch)
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_82 (Dense) | (None, 512) | 401920 |
| batch_normalization_41 (Batch Normalization) | (None, 512) | 2048 |
| dense_83 (Dense) | (None, 128) | 65664 |
| batch_normalization_42 (Batch Normalization) | (None, 128) | 512 |
| dense_84 (Dense) | (None, 10) | 1290 |

Total params: 471,434

Trainable params: 470,154

Non-trainable params: 1,280

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 276us/step - loss: 0.1825 - acc: 0.9455 - val_loss: 0.0705 - val_acc: 0.9786

Epoch 2/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0705 - acc: 0.9786 - val_loss: 0.0453 - val_acc: 0.9863

Epoch 3/20

60000/60000 [=====] - 4s 70us/step - loss: 0.0453 - acc: 0.9863 - val_loss: 0.0321 - val_acc: 0.9903

Epoch 4/20

60000/60000 [=====] - 4s 71us/step - loss: 0.0321 - acc: 0.9903 - val_loss: 0.0265 - val_acc: 0.9918

Epoch 5/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0265 - acc: 0.9918 - val_loss: 0.0213 - val_acc: 0.9930

Epoch 6/20

60000/60000 [=====] - 4s 69us/step - loss: 0.0213 - acc: 0.9930 - val_loss: 0.0204 - val_acc: 0.9933

Epoch 7/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0204 - acc: 0.9933 - val_loss: 0.0170 - val_acc: 0.9948

Epoch 8/20

60000/60000 [=====] - 4s 68us/step - loss: 0.0170 - acc: 0.9948 - val_loss: 0.0128 - val_acc: 0.9959

Epoch 9/20

60000/60000 [=====] - 4s 67us/step - loss: 0.0128 - acc: 0.9959 - val_loss: 0.0119 - val_acc: 0.9964

Epoch 10/20


60000/60000 [=====] - 4s 66us/step - loss: 0.0119 - acc: 0.9964 - val_loss: 0.0131 - val_acc: 0.9958


Epoch 11/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0131 - acc: 0.9958 - val_loss: 0.0131 - val_acc: 0.9958

Epoch 12/20

60000/60000 [=====] - 4s 66us/step - loss: 0.0131 - acc: 0.9960 - val.
Epoch 13/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0093 - acc: 0.9970 - val.
Epoch 14/20
60000/60000 [=====] - 4s 71us/step - loss: 0.0113 - acc: 0.9965 - val.
Epoch 15/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0095 - acc: 0.9966 - val.
Epoch 16/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0080 - acc: 0.9974 - val.
Epoch 17/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0072 - acc: 0.9976 - val.
Epoch 18/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0057 - acc: 0.9982 - val.
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0079 - acc: 0.9973 - val.
Epoch 20/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0093 - acc: 0.9969 - val.
Test score: 0.09142072524498654
Test accuracy: 0.9795





3.1.3 Model 1: M1 + Batch-Normalization + Dropout

In [70]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu',
                    input_shape=(input_dim,), kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(128, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

history = trainModel(model=model_drop)
plotGraph(model=model_drop, history=history)
plotWeightM1(model=model_drop)
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_85 (Dense) | (None, 512) | 401920 |
| batch_normalization_43 (Batch Normalization) | (None, 512) | 2048 |
| dropout_21 (Dropout) | (None, 512) | 0 |
| dense_86 (Dense) | (None, 128) | 65664 |
| batch_normalization_44 (Batch Normalization) | (None, 128) | 512 |
| dropout_22 (Dropout) | (None, 128) | 0 |
| dense_87 (Dense) | (None, 10) | 1290 |

Total params: 471,434

Trainable params: 470,154

Non-trainable params: 1,280

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 281us/step - loss: 0.4242 - acc: 0.8725 - val_loss: 0.3750

Epoch 2/20

60000/60000 [=====] - 4s 72us/step - loss: 0.2088 - acc: 0.9374 - val_loss: 0.2500

Epoch 3/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1627 - acc: 0.9511 - val_loss: 0.2500

Epoch 4/20

60000/60000 [=====] - 4s 73us/step - loss: 0.1401 - acc: 0.9576 - val_loss: 0.2500

Epoch 5/20

60000/60000 [=====] - 4s 71us/step - loss: 0.1227 - acc: 0.9627 - val_loss: 0.2500

Epoch 6/20

60000/60000 [=====] - 4s 72us/step - loss: 0.1150 - acc: 0.9650 - val_loss: 0.2500

Epoch 7/20

60000/60000 [=====] - 4s 72us/step - loss: 0.1033 - acc: 0.9684 - val_loss: 0.2500

Epoch 8/20

60000/60000 [=====] - 4s 72us/step - loss: 0.0935 - acc: 0.9709 - val_loss: 0.2500

Epoch 9/20

60000/60000 [=====] - 4s 75us/step - loss: 0.0895 - acc: 0.9722 - val_loss: 0.2500

Epoch 10/20

60000/60000 [=====] - 5s 76us/step - loss: 0.0828 - acc: 0.9738 - val_loss: 0.2500

Epoch 11/20

60000/60000 [=====] - 5s 77us/step - loss: 0.0788 - acc: 0.9746 - val_loss: 0.2500


Epoch 12/20


60000/60000 [=====] - 5s 76us/step - loss: 0.0745 - acc: 0.9763 - val_loss: 0.2500

Epoch 13/20

60000/60000 [=====] - 4s 74us/step - loss: 0.0704 - acc: 0.9782 - val_loss: 0.2500

Epoch 14/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0690 - acc: 0.9783 - val.
Epoch 15/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0669 - acc: 0.9793 - val.
Epoch 16/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0617 - acc: 0.9808 - val.
Epoch 17/20
60000/60000 [=====] - 4s 75us/step - loss: 0.0621 - acc: 0.9807 - val.
Epoch 18/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0564 - acc: 0.9817 - val.
Epoch 19/20
60000/60000 [=====] - 4s 73us/step - loss: 0.0575 - acc: 0.9817 - val.
Epoch 20/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0531 - acc: 0.9830 - val.
Test score: 0.054348885305505246
Test accuracy: 0.9829





3.2 Model 2 (3-layered MLP Architecture):

3.2.1 Plot Weights: Common Function

```
In [71]: def plotWeightM2(model):
    w_after = model.get_weights()

    h1_w = w_after[0].flatten().reshape(-1,1)
    h2_w = w_after[2].flatten().reshape(-1,1)
    h3_w = w_after[4].flatten().reshape(-1,1)
    out_w = w_after[6].flatten().reshape(-1,1)

    fig = plt.figure()
    plt.title("Weight matrices after model trained")
    plt.subplot(1, 4, 1)
    plt.title("Trained Wt")
    ax = sns.violinplot(y=h1_w,color='b')
    plt.xlabel('Hidden Layer 1')

    plt.subplot(1, 4, 2)
    plt.title("Trained Wt")
    ax = sns.violinplot(y=h2_w, color='r')
    plt.xlabel('Hidden Layer 2 ')
```

```

plt.subplot(1, 4, 3)
plt.title("Trained Wt")
ax = sns.violinplot(y=h3_w, color='g')
plt.xlabel('Hidden Layer 3 ')

plt.subplot(1, 4, 4)
plt.title("Trained Wt")
ax = sns.violinplot(y=out_w,color='m')
plt.xlabel('Output Layer ')
plt.show()

```

3.2.2 Input(784)-ReLu(512)-ReLu(256)-ReLu(64)-Softmax(output)

```

In [72]: model_relu = Sequential()
        model_relu.add(Dense(512, activation='relu',
                             input_shape=(input_dim,), kernel_initializer='he_normal'))
        model_relu.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(64, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(output_dim, activation='softmax'))

        print(model_relu.summary())

        history = trainModel(model=model_relu)
        plotGraph(model=model_relu, history=history)
        plotWeightM2(model=model_relu)

```

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_88 (Dense) | (None, 512) | 401920 |
| dense_89 (Dense) | (None, 256) | 131328 |
| dense_90 (Dense) | (None, 64) | 16448 |
| dense_91 (Dense) | (None, 10) | 650 |

```

Total params: 550,346
Trainable params: 550,346
Non-trainable params: 0

```

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 16s 269us/step - loss: 0.2229 - acc: 0.9336 - val

Epoch 2/20

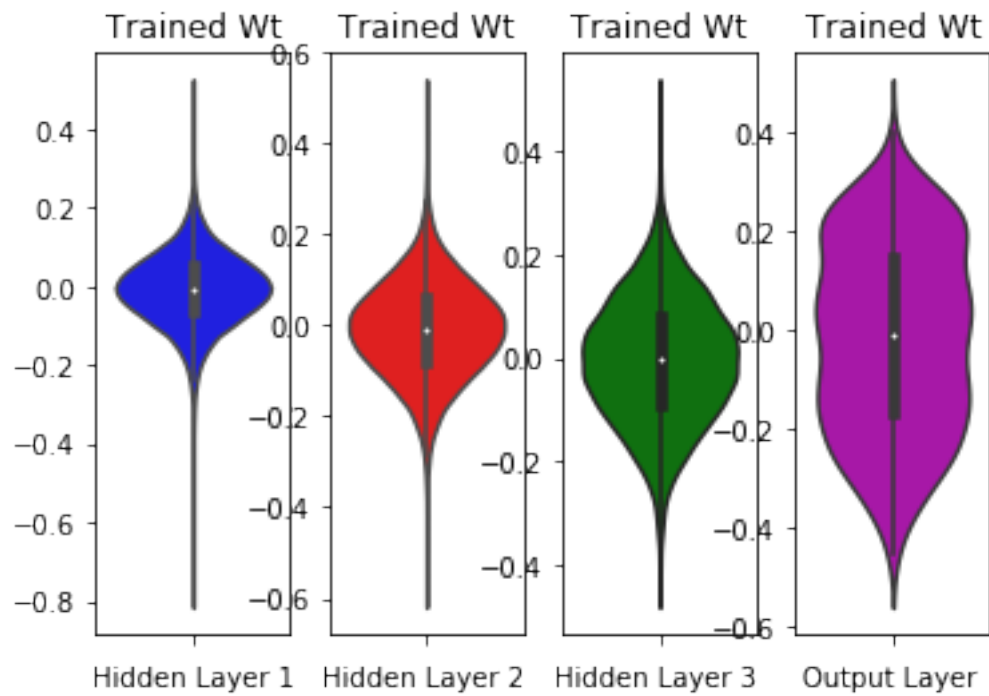

60000/60000 [=====] - 4s 66us/step - loss: 0.0834 - acc: 0.9749 - val

Epoch 3/20

```

60000/60000 [=====] - 4s 67us/step - loss: 0.0533 - acc: 0.9836 - val.
Epoch 4/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0386 - acc: 0.9874 - val.
Epoch 5/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0287 - acc: 0.9906 - val.
Epoch 6/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0282 - acc: 0.9912 - val.
Epoch 7/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0199 - acc: 0.9931 - val.
Epoch 8/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0180 - acc: 0.9942 - val.
Epoch 9/20
60000/60000 [=====] - 4s 69us/step - loss: 0.0181 - acc: 0.9940 - val.
Epoch 10/20
60000/60000 [=====] - 4s 67us/step - loss: 0.0134 - acc: 0.9957 - val.
Epoch 11/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0158 - acc: 0.9948 - val.
Epoch 12/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0121 - acc: 0.9961 - val.
Epoch 13/20
60000/60000 [=====] - 4s 65us/step - loss: 0.0101 - acc: 0.9969 - val.
Epoch 14/20
60000/60000 [=====] - 4s 66us/step - loss: 0.0112 - acc: 0.9963 - val.
Epoch 15/20
60000/60000 [=====] - 4s 68us/step - loss: 0.0110 - acc: 0.9962 - val.
Epoch 16/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0108 - acc: 0.9966 - val.
Epoch 17/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0099 - acc: 0.9969 - val.
Epoch 18/20
60000/60000 [=====] - 4s 63us/step - loss: 0.0078 - acc: 0.9977 - val.
Epoch 19/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0106 - acc: 0.9967 - val.
Epoch 20/20
60000/60000 [=====] - 4s 64us/step - loss: 0.0103 - acc: 0.9968 - val.
Test score: 0.08919506263012462
Test accuracy: 0.9807

```



3.2.3 Model 2: M2 + Batch-Normalization on 3 hidden Layers

In [73]: *# Multilayer perceptron*

```
# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution N(0,) we satisfy this condition with
# h1 =>  =(2/(ni+ni+1) = 0.039  => N(0,) = N(0,0.039)
# h2 =>  =(2/(ni+ni+1) = 0.055  => N(0,) = N(0,0.055)
# h1 =>  =(2/(ni+ni+1) = 0.120  => N(0,) = N(0,0.120)

from keras.layers.normalization import BatchNormalization

model_batch = Sequential()

model_batch.add(Dense(512, activation='relu',
                      input_shape=(input_dim,), kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())

model_batch.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())

model_batch.add(Dense(64, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()



history = trainModel(model=model_batch)
plotGraph(model=model_batch, history=history)
plotWeightM2(model=model_batch)
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_92 (Dense) | (None, 512) | 401920 |
| batch_normalization_45 (Batch Normalization) | (None, 512) | 2048 |
| dense_93 (Dense) | (None, 256) | 131328 |
| batch_normalization_46 (Batch Normalization) | (None, 256) | 1024 |
| dense_94 (Dense) | (None, 64) | 16448 |
| batch_normalization_47 (Batch Normalization) | (None, 64) | 256 |
| dense_95 (Dense) | (None, 10) | 650 |

Total params: 553,674
Trainable params: 552,010
Non-trainable params: 1,664

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 19s 313us/step - loss: 0.1918 - acc: 0.9438 - val.
Epoch 2/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0727 - acc: 0.9782 - val.
Epoch 3/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0498 - acc: 0.9850 - val.
Epoch 4/20
60000/60000 [=====] - 5s 83us/step - loss: 0.0365 - acc: 0.9885 - val.
Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0302 - acc: 0.9900 - val.
Epoch 6/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0232 - acc: 0.9928 - val.
Epoch 7/20
60000/60000 [=====] - 5s 79us/step - loss: 0.0214 - acc: 0.9931 - val.
Epoch 8/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0194 - acc: 0.9936 - val.
Epoch 9/20
60000/60000 [=====] - 6s 103us/step - loss: 0.0198 - acc: 0.9936 - val.
Epoch 10/20
60000/60000 [=====] - 6s 93us/step - loss: 0.0156 - acc: 0.9952 - val.
Epoch 11/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0136 - acc: 0.9955 - val.
Epoch 12/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0177 - acc: 0.9940 - val.
Epoch 13/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0136 - acc: 0.9954 - val.
Epoch 14/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0101 - acc: 0.9969 - val.
Epoch 15/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0094 - acc: 0.9970 - val.
Epoch 16/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0111 - acc: 0.9964 - val.
Epoch 17/20
60000/60000 [=====] - 5s 80us/step - loss: 0.0132 - acc: 0.9955 - val.
Epoch 18/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0092 - acc: 0.9969 - val.
Epoch 19/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0076 - acc: 0.9977 - val.
Epoch 20/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0081 - acc: 0.9973 - val.
Test score: 0.0944686686351095
Test accuracy: 0.9776



3.2.4 Model 2: M2 + Batch-Normalization + Dropout

In [74]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu',
                    input_shape=(input_dim,), kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(64, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

history = trainModel(model=model_drop)
plotGraph(model=model_drop, history=history)
plotWeightM2(model=model_drop)
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_96 (Dense) | (None, 512) | 401920 |
| batch_normalization_48 (Batch Normalization) | (None, 512) | 2048 |
| dropout_23 (Dropout) | (None, 512) | 0 |
| dense_97 (Dense) | (None, 256) | 131328 |
| batch_normalization_49 (Batch Normalization) | (None, 256) | 1024 |
| dropout_24 (Dropout) | (None, 256) | 0 |
| dense_98 (Dense) | (None, 64) | 16448 |
| batch_normalization_50 (Batch Normalization) | (None, 64) | 256 |
| dropout_25 (Dropout) | (None, 64) | 0 |

```

-----
dense_99 (Dense)                (None, 10)                650
=====

```

```

Total params: 553,674
Trainable params: 552,010
Non-trainable params: 1,664

```

```

-----
Train on 60000 samples, validate on 10000 samples



```

```

Epoch 1/20
60000/60000 [=====] - 19s 310us/step - loss: 0.6132 - acc: 0.8129 - val.
Epoch 2/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2699 - acc: 0.9229 - val.
Epoch 3/20
60000/60000 [=====] - 5s 89us/step - loss: 0.2023 - acc: 0.9429 - val.
Epoch 4/20
60000/60000 [=====] - 5s 88us/step - loss: 0.1724 - acc: 0.9514 - val.
Epoch 5/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1530 - acc: 0.9574 - val.
Epoch 6/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1337 - acc: 0.9614 - val.
Epoch 7/20
60000/60000 [=====] - 5s 86us/step - loss: 0.1246 - acc: 0.9656 - val.
Epoch 8/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1153 - acc: 0.9673 - val.
Epoch 9/20
60000/60000 [=====] - 5s 87us/step - loss: 0.1118 - acc: 0.9674 - val.
Epoch 10/20
60000/60000 [=====] - 5s 90us/step - loss: 0.1024 - acc: 0.9708 - val.
Epoch 11/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0969 - acc: 0.9715 - val.
Epoch 12/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0904 - acc: 0.9744 - val.
Epoch 13/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0866 - acc: 0.9748 - val.
Epoch 14/20
60000/60000 [=====] - 5s 88us/step - loss: 0.0820 - acc: 0.9759 - val.
Epoch 15/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0806 - acc: 0.9767 - val.
Epoch 16/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0777 - acc: 0.9769 - val.
Epoch 17/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0745 - acc: 0.9775 - val.
Epoch 18/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0712 - acc: 0.9787 - val.
Epoch 19/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0706 - acc: 0.9792 - val.
Epoch 20/20
60000/60000 [=====] - 5s 90us/step - loss: 0.0672 - acc: 0.9801 - val.

```

Test score: 0.06557027021200047
Test accuracy: 0.9828



3.3 Model 3 (5-layered MLP Architecture):

3.3.1 Plot Weights: Common Function

```
In [75]: def plotWeightM3(model):
    w_after = model.get_weights()

    h1_w = w_after[0].flatten().reshape(-1,1)
    h2_w = w_after[2].flatten().reshape(-1,1)
    h3_w = w_after[4].flatten().reshape(-1,1)
    h4_w = w_after[6].flatten().reshape(-1,1)
    h5_w = w_after[8].flatten().reshape(-1,1)
    out_w = w_after[10].flatten().reshape(-1,1)

    fig = plt.figure()
    plt.title("Weight matrices after model trained")
    plt.subplot(1, 6, 1)
    plt.title("Weights")
    ax = sns.violinplot(y=h1_w,color='b')
    plt.xlabel('Hidden 1')

    plt.subplot(1, 6, 2)
    plt.title("Weights")
    ax = sns.violinplot(y=h2_w, color='r')
    plt.xlabel('Hidden 2 ')

    plt.subplot(1, 6, 3)
    plt.title("Weights")
    ax = sns.violinplot(y=h3_w, color='g')
    plt.xlabel('Hidden 3 ')

    plt.subplot(1, 6, 4)
    plt.title("Weights")
    ax = sns.violinplot(y=h4_w, color='c')
    plt.xlabel('Hidden 4 ')

    plt.subplot(1, 6, 5)
    plt.title("Weights")
    ax = sns.violinplot(y=h5_w, color='m')
    plt.xlabel('Hidden 5 ')

    plt.subplot(1, 6, 6)
    plt.title("Weights")
    ax = sns.violinplot(y=out_w,color='y')
```

```
plt.xlabel('Out Layer')
plt.show()
```

3.3.2 Input(784)-ReLu(512)-ReLu(256)-ReLu(144)-ReLu(96)-ReLu(36)-Softmax(output)

```
In [76]: model_relu = Sequential()
        model_relu.add(Dense(512, activation='relu',
                             input_shape=(input_dim,), kernel_initializer='he_normal'))
        model_relu.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(144, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(96, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(36, activation='relu', kernel_initializer='he_normal'))
        model_relu.add(Dense(output_dim, activation='softmax'))

        print(model_relu.summary())

        history = trainModel(model=model_relu)
        plotGraph(model=model_relu, history=history)
        plotWeightM3(model=model_relu)
```

| Layer (type) | Output Shape | Param # |
|-------------------|--------------|---------|
| dense_100 (Dense) | (None, 512) | 401920 |
| dense_101 (Dense) | (None, 256) | 131328 |
| dense_102 (Dense) | (None, 144) | 37008 |
| dense_103 (Dense) | (None, 96) | 13920 |
| dense_104 (Dense) | (None, 36) | 3492 |
| dense_105 (Dense) | (None, 10) | 370 |

Total params: 588,038

Trainable params: 588,038

Non-trainable params: 0

None

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 17s 286us/step - loss: 0.2376 - acc: 0.9294 - val

Epoch 2/20

60000/60000 [=====] - 5s 81us/step - loss: 0.0914 - acc: 0.9722 - val

Epoch 3/20



60000/60000 [=====] - 5s 77us/step - loss: 0.0592 - acc: 0.9814 - val

Epoch 4/20

```

60000/60000 [=====] - 5s 77us/step - loss: 0.0449 - acc: 0.9856 - val.
Epoch 5/20
60000/60000 [=====] - 5s 86us/step - loss: 0.0368 - acc: 0.9880 - val.
Epoch 6/20
60000/60000 [=====] - 5s 76us/step - loss: 0.0303 - acc: 0.9900 - val.
Epoch 7/20
60000/60000 [=====] - 5s 82us/step - loss: 0.0237 - acc: 0.9921 - val.
Epoch 8/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0262 - acc: 0.9919 - val.
Epoch 9/20
60000/60000 [=====] - 5s 85us/step - loss: 0.0187 - acc: 0.9940 - val.
Epoch 10/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0204 - acc: 0.9938 - val.
Epoch 11/20
60000/60000 [=====] - 5s 81us/step - loss: 0.0204 - acc: 0.9934 - val.
Epoch 12/20
60000/60000 [=====] - 5s 87us/step - loss: 0.0174 - acc: 0.9944 - val.
Epoch 13/20
60000/60000 [=====] - 5s 91us/step - loss: 0.0138 - acc: 0.9955 - val.
Epoch 14/20
60000/60000 [=====] - 5s 77us/step - loss: 0.0129 - acc: 0.9955 - val.
Epoch 15/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0147 - acc: 0.9954 - val.
Epoch 16/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0133 - acc: 0.9960 - val.
Epoch 17/20
60000/60000 [=====] - 4s 74us/step - loss: 0.0133 - acc: 0.9960 - val.
Epoch 18/20
60000/60000 [=====] - 4s 70us/step - loss: 0.0090 - acc: 0.9971 - val.
Epoch 19/20
60000/60000 [=====] - 4s 72us/step - loss: 0.0140 - acc: 0.9960 - val.
Epoch 20/20
60000/60000 [=====] - 5s 78us/step - loss: 0.0094 - acc: 0.9972 - val.
Test score: 0.09236639852523267
Test accuracy: 0.9818

```



3.3.3 Model 3: M3 + Batch-Normalization on 5 hidden Layers

In [77]: *# Multilayer perceptron*

```
# https://intoli.com/blog/neural-network-initialization/
# If we sample weights from a normal distribution N(0,) we satisfy this condition with
# h1 =>  =(2/(ni+ni+1)) = 0.039  => N(0,) = N(0,0.039)
# h2 =>  =(2/(ni+ni+1)) = 0.055  => N(0,) = N(0,0.055)
# h1 =>  =(2/(ni+ni+1)) = 0.120  => N(0,) = N(0,0.120)
```

```
from keras.layers.normalization import BatchNormalization
```

```
model_batch = Sequential()
```

```
model_batch.add(Dense(512, activation='relu',
                      input_shape=(input_dim,), kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(144, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(96, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(36, activation='relu', kernel_initializer='he_normal'))
model_batch.add(BatchNormalization())
```

```
model_batch.add(Dense(output_dim, activation='softmax'))
```

```
model_batch.summary()
```

```
history = trainModel(model=model_batch)
plotGraph(model=model_batch, history=history)
plotWeightM3(model=model_batch)
```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_106 (Dense) | (None, 512) | 401920 |
| batch_normalization_51 (Batch Normalization) | (None, 512) | 2048 |
| dense_107 (Dense) | (None, 256) | 131328 |
| batch_normalization_52 (Batch Normalization) | (None, 256) | 1024 |

| | | |
|--|-------------|-------|
| dense_108 (Dense) | (None, 144) | 37008 |
| batch_normalization_53 (Batch Normalization) | (None, 144) | 576 |
| dense_109 (Dense) | (None, 96) | 13920 |
| batch_normalization_54 (Batch Normalization) | (None, 96) | 384 |
| dense_110 (Dense) | (None, 36) | 3492 |
| batch_normalization_55 (Batch Normalization) | (None, 36) | 144 |
| dense_111 (Dense) | (None, 10) | 370 |

Total params: 592,214
 Trainable params: 590,126
 Non-trainable params: 2,088

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 19s 325us/step - loss: 0.2446 - acc: 0.9314 - val_loss: 0.1000

Epoch 2/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0890 - acc: 0.9734 - val_loss: 0.0400

Epoch 3/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0637 - acc: 0.9801 - val_loss: 0.0300

Epoch 4/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0492 - acc: 0.9849 - val_loss: 0.0200

Epoch 5/20

60000/60000 [=====] - 6s 96us/step - loss: 0.0414 - acc: 0.9865 - val_loss: 0.0200

Epoch 6/20

60000/60000 [=====] - 6s 99us/step - loss: 0.0334 - acc: 0.9895 - val_loss: 0.0200

Epoch 7/20

60000/60000 [=====] - 6s 97us/step - loss: 0.0285 - acc: 0.9904 - val_loss: 0.0200

Epoch 8/20

60000/60000 [=====] - 6s 107us/step - loss: 0.0282 - acc: 0.9903 - val_loss: 0.0200

Epoch 9/20

60000/60000 [=====] - 6s 101us/step - loss: 0.0276 - acc: 0.9908 - val_loss: 0.0200

Epoch 10/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0237 - acc: 0.9923 - val_loss: 0.0200

Epoch 11/20

60000/60000 [=====] - 6s 98us/step - loss: 0.0202 - acc: 0.9936 - val_loss: 0.0200

Epoch 12/20

60000/60000 [=====] - 6s 95us/step - loss: 0.0183 - acc: 0.9938 - val_loss: 0.0200


Epoch 13/20


60000/60000 [=====] - 6s 101us/step - loss: 0.0174 - acc: 0.9943 - val_loss: 0.0200

Epoch 14/20

60000/60000 [=====] - 6s 104us/step - loss: 0.0192 - acc: 0.9937 - val_loss: 0.0200

Epoch 15/20
60000/60000 [=====] - 6s 107us/step - loss: 0.0128 - acc: 0.9959 - val.
Epoch 16/20
60000/60000 [=====] - 6s 108us/step - loss: 0.0142 - acc: 0.9956 - val.
Epoch 17/20
60000/60000 [=====] - 7s 109us/step - loss: 0.0137 - acc: 0.9957 - val.
Epoch 18/20
60000/60000 [=====] - 6s 104us/step - loss: 0.0149 - acc: 0.9951 - val.
Epoch 19/20
60000/60000 [=====] - 6s 94us/step - loss: 0.0137 - acc: 0.9956 - val.
Epoch 20/20
60000/60000 [=====] - 6s 98us/step - loss: 0.0118 - acc: 0.9961 - val.
Test score: 0.07374950621149037
Test accuracy: 0.9813





3.3.4 Model 3: M3 + Batch-Normalization + Dropout

In [78]: # <https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization>

```
from keras.layers import Dropout

model_drop = Sequential()

model_drop.add(Dense(512, activation='relu',
                    input_shape=(input_dim,), kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(256, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(144, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(96, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))
```

```

model_drop.add(Dense(36, activation='relu', kernel_initializer='he_normal'))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.5))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

history = trainModel(model=model_drop)
plotGraph(model=model_drop, history=history)
plotWeightM3(model=model_drop)



```

| Layer (type) | Output Shape | Param # |
|--|--------------|---------|
| dense_112 (Dense) | (None, 512) | 401920 |
| batch_normalization_56 (Batch Normalization) | (None, 512) | 2048 |
| dropout_26 (Dropout) | (None, 512) | 0 |
| dense_113 (Dense) | (None, 256) | 131328 |
| batch_normalization_57 (Batch Normalization) | (None, 256) | 1024 |
| dropout_27 (Dropout) | (None, 256) | 0 |
| dense_114 (Dense) | (None, 144) | 37008 |
| batch_normalization_58 (Batch Normalization) | (None, 144) | 576 |
| dropout_28 (Dropout) | (None, 144) | 0 |
| dense_115 (Dense) | (None, 96) | 13920 |
| batch_normalization_59 (Batch Normalization) | (None, 96) | 384 |
| dropout_29 (Dropout) | (None, 96) | 0 |
| dense_116 (Dense) | (None, 36) | 3492 |
| batch_normalization_60 (Batch Normalization) | (None, 36) | 144 |
| dropout_30 (Dropout) | (None, 36) | 0 |
| dense_117 (Dense) | (None, 10) | 370 |

Total params: 592,214
Trainable params: 590,126
Non-trainable params: 2,088

Train on 60000 samples, validate on 10000 samples

Epoch 1/20
60000/60000 [=====] - 20s 341us/step - loss: 1.3673 - acc: 0.5682 - va
Epoch 2/20
60000/60000 [=====] - 7s 109us/step - loss: 0.5358 - acc: 0.8528 - va
Epoch 3/20
60000/60000 [=====] - 6s 108us/step - loss: 0.3650 - acc: 0.9085 - va
Epoch 4/20
60000/60000 [=====] - 6s 108us/step - loss: 0.3002 - acc: 0.9260 - va
Epoch 5/20
60000/60000 [=====] - 7s 109us/step - loss: 0.2589 - acc: 0.9363 - va
Epoch 6/20
60000/60000 [=====] - 6s 108us/step - loss: 0.2342 - acc: 0.9440 - va
Epoch 7/20
60000/60000 [=====] - 6s 108us/step - loss: 0.2094 - acc: 0.9494 - va
Epoch 8/20
60000/60000 [=====] - 6s 107us/step - loss: 0.1986 - acc: 0.9531 - va
Epoch 9/20
60000/60000 [=====] - 7s 110us/step - loss: 0.1835 - acc: 0.9567 - va
Epoch 10/20
60000/60000 [=====] - 6s 108us/step - loss: 0.1738 - acc: 0.9589 - va
Epoch 11/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1660 - acc: 0.9616 - va
Epoch 12/20
60000/60000 [=====] - 7s 110us/step - loss: 0.1564 - acc: 0.9627 - va
Epoch 13/20
60000/60000 [=====] - 7s 110us/step - loss: 0.1461 - acc: 0.9659 - va
Epoch 14/20
60000/60000 [=====] - 6s 107us/step - loss: 0.1475 - acc: 0.9655 - va
Epoch 15/20
60000/60000 [=====] - 6s 106us/step - loss: 0.1407 - acc: 0.9672 - va
Epoch 16/20
60000/60000 [=====] - 7s 108us/step - loss: 0.1378 - acc: 0.9675 - va
Epoch 17/20
60000/60000 [=====] - 7s 111us/step - loss: 0.1303 - acc: 0.9687 - va
Epoch 18/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1238 - acc: 0.9712 - va
Epoch 19/20
60000/60000 [=====] - 7s 123us/step - loss: 0.1223 - acc: 0.9719 - va
Epoch 20/20
60000/60000 [=====] - 7s 116us/step - loss: 0.1145 - acc: 0.9733 - va
Test score: 0.0764596716644708
Test accuracy: 0.9814



4 Summary Statistics

| Model | Total Parameters | Accuracy | Loss vs Epoch Plot |
|---|------------------|--------------|--------------------|
| Model 1 (784x512x128x10) | 468,874 | 98.18 | Diverging |
| M1 + Batch-Normalization | 471,434 | 98.06 | Diverging |
| M1 + Batch-Normalization+Dropout | 471,434 | 98.32 | Converging |
| Model 2 (784x512x256x64x10) | 550,346 | 97.85 | Diverging |
| M2 + Batch-Normalization | 553,674 | 98.06 | Diverging |
| M2 + Batch-Normalization+Dropout | 553,674 | 98.24 | Converging |
| Model 3 (784x512x256x144x96x10) | 588,038 | 97.96 | Diverging |
| M3 + Batch-Normalization | 592,214 | 97.88 | Diverging |
| M3 + Batch-Normalization+Dropout | 592,214 | 98.26 | Converging |

5 Conclusions

1. The **difference in accuracy between 2, 3 & 5 layered networks is very small**. This could be due to the simplicity and small size of input data.
2. The '**cross entropy loss vs epoch**' plot for train and test data is found **diverging, when the dropout layer is not added**. This means reduction in training loss but increase in test loss at the same time, **indicative of overfitting**.
3. Thus, **addition of dropout layer is found as a good regularization** in practice.
4. The accuracy is also more when Batch Normalization and dropout layers are added.
5. All distributions of trained weights along all the layers are expectedly found as Gaussian curves.
6. For MNIST problem, **model M1 coupled with Batch Normalization and Dropout seems to be the best bet**.